

# Machine Learning in Security

---

CS463/ECE424

University of Illinois



## Definitions

Spam Classification using Logistic Regression

Anomaly Detection through Deep Learning

Challenges for Machine Learning in Security



# What is Machine Learning?

## Traditional Programming


When we know how to do things




```
int addition (int a, int b)
{
    int r;
    r = a + b;
    return r;
}
```

# What is Machine Learning?


## Traditional Programming



## Machine Learning



When we don't know how to do it, but we have some examples



Cat??

# What is Machine Learning?

---

A computer program is said to learn from **experience E** with respect to some class of **tasks T** and **performance measure P**, if its performance at tasks in T, as measured by P, improves with experience E.

-- Tom Mitchell, *Machine Learning*

# Steps towards Designing a ML System

---

- Step 1: Choosing the Training Experience (i.e., training dataset)
- Step 2: Modeling the Transformation
- Step 3: Choosing the Input & Output Representations
- Step 4: Choosing a Transformation Function Approximation
- Step 5: Evaluation

# When To Use Machine Learning?

---

- **When patterns exist in the data**
  - Even if we don't know what they are
- **When we cannot pin down the functional relationships mathematically (in-closed form)**
  - Else we would just code up the algorithm
- **When we have lots of data**
  - Labeled training sets are harder to come by than unlabeled data
  - Data is of high-dimension
- **When we want to discover lower-dimension representations**



# Example: Spam Filtering



- **Task T:** classifying emails into two categories (spam, ham)
- **Performance measure P:** percent of emails correctly classified
- **Training Experience E:** a database of emails


# Step 1: Choosing the Training Dataset

---

- Training Dataset:
  - A database of emails
- What feedback can be provided to the learner?
  - A database of labeled emails
- How well does the **training dataset** represent the distribution of examples over which the **final system** performance  $P$  must be measured?
  - A database of labeled emails that represent the distribution of all the emails

# Step 1: Choosing the Training Experience

---



Cats v.s. Dogs

# Step 2: Modeling the Transformation

---

- **Task T:** classifying emails into 2 categories (Spam, Ham)
- Transformation (target) function  $V: A \rightarrow B$ 
  - What's in the training examples?
    - A: Email contents (a collection of words)
  - What should be the output?
    - B: {Spam (1) , Ham (0)}




# Step 3: Choosing the Input & Output Representations

---

- How do we represent the model inputs and outputs?
  - Inputs could be categorical, numerical, binary, sequential ...
  - We use code + math; they need “numbers”
- **Feature generation**
  - Abstract  $V$ : Email contents  $\rightarrow \{0,1\}$
  - Realized  $V'$ :  $\mathbf{x} = (x_1, x_2, \dots, x_n) \mapsto y \in \{0,1\}$ 
    - $x_i \in \{0,1\}$  represents whether a word  $w_i$  is in the email
- **Feature selection**
  - To simplify the model (save time, avoid overfitting...)

# Step 4: Choosing the Transformation Function Approximation



# Logistic Regression (1)


---

- $V': \mathbf{x} = (x_1, x_2, \dots, x_n) \mapsto y \in \{0,1\}$
- How to design  $V'$ ?
  - **Step 1:** Combine  $x_1, x_2, \dots, x_n$  to get a “spaminess” value
  - **Step 2:** Convert the “spaminess” value into a probability  $P(\text{Spam})$ 
    - Conversion enables the classification use-case; else it is useful for regression
  - **Step 3:** Make predictions on  $y$  based on  $P(\text{Spam})$ 
    - e.g.,  $y = 1$  when  $P(\text{Spam}) > 0.5$

# Logistic Regression (2)

---

- Step 1: Combine  $x_1, x_2, \dots, x_n$  to get a “spaminess” value
  - Assume the existence of a weight vector  $\boldsymbol{\theta} = (\theta_1, \theta_2, \dots, \theta_n)$
  - We define spaminess as the linear transformation  $\boldsymbol{\theta}^T \cdot \mathbf{x}$




# Logistic Regression (3)

---

- Step 2: Convert the “spaminess” value into a probability  $P(\text{Spam})$ 
  - Logistic function

$$h_{\theta}(x) = g(\theta^T X) = \frac{1}{1 + e^{-\theta^T x}}$$




Plot of logistic function  $g$

# Logistic Regression (4)

---

- Step 3: Make predictions on  $y$  based on  $P(\text{Spam})$



# Logistic Regression: Training

---

- How do we determine the “best” value of  $\theta$ ?
  - For a given  $\theta$  and some labeled examples, how do we know whether  $\theta$  is good enough? i.e., is the best predictor of spam vs. ham?
- **Define a loss function (e.g., log loss)**

$$\text{Log Loss} = \sum_{(x,y) \in D} -y \log(y') - (1 - y) \log(1 - y')$$

where:

- $(x, y) \in D$  is the dataset containing many labeled examples, which are  $(x, y)$  pairs.
- $y$  is the label in a labeled example. Since this is logistic regression, every value of  $y$  must either be 0 or 1.
- $y'$  is your model's prediction (somewhere between 0 and 1), given the set of features in  $x$ .

# Logistic Regression: Training

---

- How do we determine the “best” value of  $\theta$ ?
  - For a given  $\theta$  and some labeled examples, how do we know whether  $\theta$  is good enough? i.e., is the best predictor of spam vs. ham?
- **Define a loss function** (e.g., log loss)
  - Wrong predictions -> large loss
  - Correct predictions -> small loss
- Run **optimization algorithms** to find  $\theta$ , minimize the loss
  - e.g., Stochastic Gradient Descent (SGD)

# Step 5: Evaluation

---

- Ground Truth
  - V: Email contents  $\rightarrow \{0,1\}$
- Hold out Method
  - Randomly partitioned data into two independent sets: a **test set**, a **training set**
  - Use **test set** instead of **training set** when assessing accuracy
- Cross-validation (k-fold)
  - Randomly partition the data into  $k$  **mutually exclusive** subsets, each approximately equal size
  - At  $i$ -th iteration, use  $D_i$  as **test set** and **others as training set**


Test Set

Training Set

# Step 5: Evaluation

---

- Overfitting:



# Step 5: Evaluation


---



Cats v.s. Dogs

# Step 5: Evaluation

- Overfitting:



- Confusion Matrix:

	Predicted Spam	Predicted Ham
Spam	True Positive (TP)	False Negative (FN)
Ham	False Positive (FP)	True Negative (TN)

# Summary: Designing a ML System

---

- Step 1: Choosing the Training Experience (i.e., training dataset)
- Step 2: Modeling the Transformation
- Step 3: Choosing the Input & Output Representations
- Step 4: Choosing a Transformation Function Approximation
- Step 5: Evaluation

# DeepLog: Anomaly Detection through Deep Learning

---

- Anomaly Detection from System Logs
  - Identify **abnormal system behavior** from **large volume of system logs**
- Challenges
  - Large volume of data
  - Sequential data
  - Unstructured data
- Why deep learning?
  - Widely used for natural language processing (NLP)
  - *Log can be viewed as **a structured language!***

# Step 1: Choosing the Training Data (1)

---

- What data do we have?
  - Large volume of log entries from normal system execution path
  - A few log entries of known attacks



# Step 1: Choosing the Training Data (2)

---

- What data should we use?
  - Training: **normal logs**
  - Testing: **normal logs** and **attack logs**
- **Advantages:**
  - Prevent overfitting
  - Test the system's behavior on unseen attacks
- **Disadvantages:**
  - May classify *any unseen behaviors as attacks* (i.e., false positives)

# Step 2: Modeling the Transformation

---

- **Outputs:** normal (-) v.s. abnormal (+)
- **Inputs:** Log entries from OpenStack VM deletion task (unstructured)
  - *t1* Deletion of file1 complete
  - *t2* Took 0.61 seconds to deallocate network ...
  - *t3* VM Stopped (Lifecycle Event)
- **Structured representation:**
  - Log key
  - **Parameter value** (e.g., t1, file1)

## Step 3: Choosing the Input & Output Representations (1)

- The total number of **distinct log keys is constant**.
  - Log keys:  $K = \{k_1, k_2, \dots, k_n\}$
  - Parameter value vectors: (time interval, other parameter values)

<b>log message (log key underlined)</b>	<b>log key</b>	<b>parameter value vector</b>
$t_1$ <u>Deletion of file1 complete</u>	$k_1$	$[t_1 - t_0, \text{file1Id}]$
$t_2$ <u>Took 0.61 seconds to deallocate network ...</u>	$k_2$	$[t_2 - t_1, 0.61]$
$t_3$ <u>VM Stopped (Lifecycle Event)</u>	$k_3$	$[t_3 - t_2]$
...	...	...

**Table 1: Log entries from OpenStack VM deletion task.**

## Step 3: Choosing the Input & Output Representations (2)


---

- Representation of Inputs:
  - **Log Keys:** structured, sequential, nominal
  - **Parameter Values:** structured, sequential, numerical (e.g., time, duration) or nominal (e.g., process id)
  - Different log keys have **different structures** for parameter values
- How to combine the inputs of different structures?
  - Train multiple models

## Step 3: Choosing the Input & Output Representations (3)

- **Model 1: Log key anomaly** detection model

- Log keys:  $K = \{k_1, k_2, \dots, k_n\}$
- Input: A window  $w$  of the  $h$  most recent log keys  $w = \{m_{t-h}, \dots, m_{t-2}, m_{t-1}\}$ , where  $m_i \in K$
- Output:  $\Pr[m_t = k_i | w]$  for each  $k_i \in K$ , ( $i = 1, \dots, n$ )



## Step 3: Choosing the Input & Output Representations (4)

---


- **Model 2: Parameter value** anomaly detection models
  - View each parameter value vector sequence (for a log key) as a separate time series
  - Train a separate model for each distinct log key value to predict the next parameter value
- Two steps of detecting anomaly
  - Predict the the next log key and parameter values
  - Compare the prediction against the observed log entry
    - Mark as anomaly if the probability for the observed log entry is low (not in the top  $g$  candidates)

## Step 3: Choosing the Input & Output Representations (5)

---

- Two steps of detecting anomaly
  - Predict the the next log key and parameter values
  - Compare the prediction against the observed log entry
    - Mark as anomaly if the probability for the observed log entry is low (not in the top  $g$  candidates)

# Step 3: Choosing the Input & Output Representations (6)



## Step 4: Choosing a Transformation Function Approximation

---

- Long Short-Term Memory (LSTM) Network
  - Has the capability of remembering previous inputs
  - Suitable for sequential data
  - A gentle walk through on LSTM networks (**optional**, 25 minutes):  
<https://www.youtube.com/watch?v=WCUNPb-5EYI>

# Step 5: Evaluation – Log Key Model (1)

---

- Hadoop-Distributed File System (HDFS) Dataset
  - System logs generated by map-reduce jobs on more than 200 Amazon’s EC2 nodes
  - Labeled by domain experts
  - Log entries are grouped into sessions

Log data set	Number of sessions		<i>n</i> : Number of log keys
	Training data (if needed)	Test data	
HDFS	4,855 normal; 1,638 abnormal	553,366 normal; 15,200 abnormal	29

- DeepLog **does not use the abnormal training data**

# Step 5: Evaluation – Log Key Model (2)


---

	Predicted as Normal	Predicted as Abnormal
Normal	552,533 (True Negative)	833 (False Positive)
Abnormal	619 (False Negative)	14581 (True Positive)

- Precision = True Positive / (True Positive + False Positive) = 94.60%
- Recall = True Positive / (True Positive + False Negative) = 95.93%

# Step 5: Evaluation – Parameter Value Model

- **OpenStack Log Dataset**
  - Run VM-related tasks
  - Inject anomalies at different execution points
- **Mean-squared error (MSE)** between the parameter value vector and the prediction output vector from DeepLog



# Challenges for Machine Learning in Security

---

- Outlier Detection
- High Cost of Errors
- Semantic Gap
- Diversity with Data
- Difficulties with Evaluations



# Case Study: Outlier Detection

---

- ML needs large number of representatives for **each** class
  - What happens when  $P(\textit{Spam})$  is very small?
- Not good at finding previously unknown malicious activities

# High Cost of Errors

- **Example:** suppose a system generates
  - 1,000,000 audit records per day;
  - 10 audit records per intrusion;
  - **Two intrusions** per day.
- Intrusion:  $I$ , Alarm:  $A$
- Detection rate:  $P(A|I) = 99.9\%$
- False alarm rate:  $P(A|\neg I) = 0.02\%$
- Given a detected record, what's the probability that the record represents a true intrusion?

$$P(I|A) = \frac{P(A|I)P(I)}{P(A|I)P(I) + P(A|\neg I)(1 - P(I))} = 9\%$$

S. Axelsson, "The Base-Rate Fallacy and Its Implications for the Difficulty of Intrusion Detection," in *Proc. ACM Conference on Computer and Communications Security*, 1999.



# Semantic Gaps


---

- Difficult to transfer results into **actionable** report for the network operator
- Difficult to find the difference between “abnormal activity” and attacks
- Not interpretable! Unclear what the system learned
  - What do false positives and false negatives mean?
  - What features are used to produce correct results?

# Diversity with Data and Concept Drift

---

- Large variability in network traffic over short time intervals



**Figure 5: Bytes Per Day**

# Difficulties with Evaluations

---

- Lack of (reliable) “ground truth”
- Outdated datasets
- Highly sensitive information (e.g., network traffic can include personal communications and business secrets)
- Difficulties with simulation and anonymization

# Reading

---

- [1] Androutsopoulos, Ion, et al. "An evaluation of naive bayesian anti-spam filtering." *arXiv preprint cs/0006013* (2000).
- [2] Du, Min, et al. "DeepLog: Anomaly Detection and Diagnosis from System Logs through Deep Learning." *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*. ACM, 2017.
- [3] Sommer, Robin, and Vern Paxson. "Outside the closed world: On using machine learning for network intrusion detection." *Security and Privacy (SP), 2010 IEEE Symposium on*. IEEE, 2010.

# Discussion Questions

---

- How can you attack the spam filtering model we discussed?
  - Can you get around the filtering and send a spam to a user's inbox?
  - Can you trick the algorithm to filter a ham email?
- Do you think ML will replace human analysts in detecting security threats? Why or why not?