

Crypto Constructs

CS463/ECE424

University of Illinois



Crypto Constructs

Homomorphic Encryption

Private Set Intersection

Searchable Encryption



Pure
Theory



Foundations for
Applicable
Techniques

Recap: Basic Crypto Concepts

- Symmetric key cryptography
 - Same key is used to encrypt and decrypt
 - Block ciphers, stream ciphers
- Public key cryptography
 - Public key for encryption, private key for decryption
 - E.g., RSA
- Collision-resistant hash functions



Background: Threat Model (1)

Attack Goal: get target plaintext

- Ciphertext-only attacks
 - Attacker only has access to the ciphertext
 - Most realistic
- Known-plaintext attacks
 - Attacker has access to a set of (plaintext, ciphertext) pairs
- Chosen-plaintext attacks
 - Attacker can pick arbitrary plaintext and obtain corresponding ciphertext
- Chosen-ciphertext attacks
 - Attacker can pick arbitrary ciphertext and obtain corresponding plaintext
 - Strongest attacker

Background: Threat Model (2)

Indistinguishability under Chosen Plaintext Attack (IND-CPA)

- Adversary can't distinguish pairs of ciphertexts with respect to their plaintexts
 - I.e., Give the attacker $C1 = \text{Enc}(P1)$, $C2 = \text{Enc}(P2)$ and $P1, P2$, and ask the attacker to create the mapping b/w P_i and C_j
- Requires **nondeterministic** encryption scheme ($E_K(m)$ is really $E_K(m, r)$ for some random r)

Why do we need randomized encryption?

(1)

IND-CPA Game:

- First the challenger generates an encryption keypair and sends the public key pk to the adversary. (It keeps the secret key.)
- Next, the adversary selects a pair of messages M_0, M_1 (of equal length) and sends them to the challenger.
- The challenger picks a random bit $b \in \{0,1\}$ and encrypts one of the two messages as $C^* \leftarrow Enc(M_b, pk)$. It sends back C^* to the adversary.
- Finally, the adversary outputs a guess b' . We say the adversary “wins” if it guesses correctly: that is, if $b' = b$.

Why do we need randomized encryption?

(2)

- First the challenger generates an encryption keypair and sends the public key pk to the adversary. (It keeps the secret key.)
- Next, the adversary selects a pair of messages M_0, M_1 (of equal length) and sends them to the challenger.
- The challenger picks a random bit $b \in \{0,1\}$ and encrypts one of the two messages as $C^* \leftarrow \text{Enc}(M_b, pk)$. It sends back C^* to the adversary.
- Finally, the adversary outputs a guess b' . We say the adversary “wins” if it guesses correctly: that is, if $b' = b$.

If encryption is not randomized..

- The adversary picks two messages M_0, M_1 and then encrypts both of them using the public key.
- When the adversary receives the ciphertext C^* , he just compares that ciphertext to the two he generated himself.
- Voila, the adversary can always figure out which message was encrypted i.e., encryption fails!!


Homomorphic Encryption



What if we could...

1. Encrypt data
2. Send it to the cloud
3. Ask the cloud to perform operations
 - Compute, search, sort


Keeping data encrypted throughout the operation!



Who would be interested in such technique?

Privacy Homomorphisms

- [RAD78] Originally idea introduced by Rivest, Adleman, and Dertouzos
- Proposed several **privacy homomorphisms**, but none of them were secure against chosen-plaintext attacks
 - Mostly because the encryption scheme is not randomized



Privacy homomorphism: Operators (\square, \circ) such that $E(x) \circ E(y) = E(x \square y)$


Homomorphic Encryption

- Fully Homomorphic Encryption (FHE)
 - Two operations: e.g., **addition** and **multiplication**
 - $E(x (y + z)) = E(x) \Delta (E(y) \circ E(z))$
 - [Gentry09] First scheme
 - Not efficient
- Partially Homomorphic Encryption (PHE)
 - Only one operation: e.g., **only multiplication**
 - $E(x y) = E(x) \Delta E(y)$
 - Many public-key cryptosystems are partially homomorphic
 - e.g., RSA - Fairly efficient

Plain RSA

Setup:

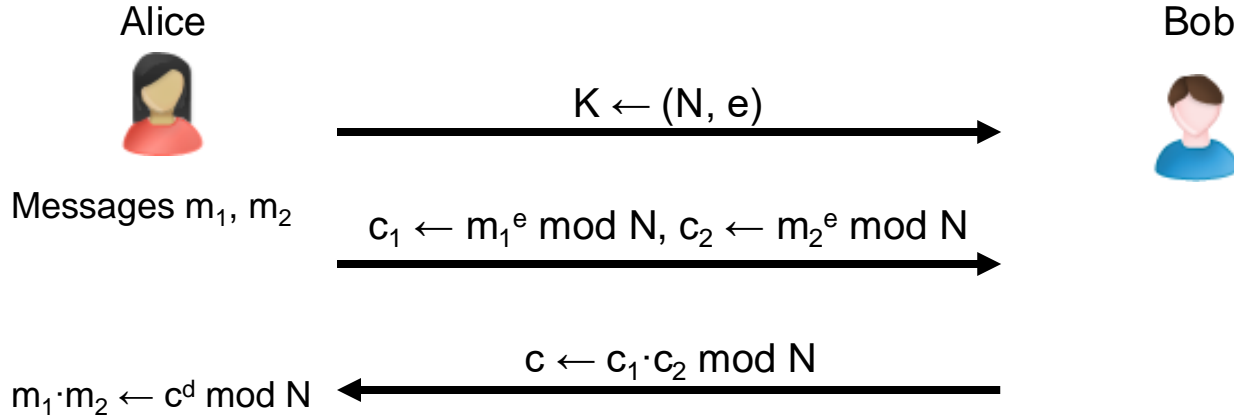
- p and q large primes, $N = pq$, $z = (p-1)(q-1)$,
- Take e co-prime with z , and calculate $d = e^{-1} \bmod z$,
- $K' = (N, d)$ is the private key
- $K = (N, e)$ is the public key



RSA

Setup:

- p and q large primes, $N = pq$, $z = (p-1)(q-1)$,
- Take e coprime with z , $d = e^{-1} \bmod z$
- $K' = (N, d)$



Plain RSA is a privacy homomorphism with respect to multiplication: $E_K(xy) = E_K(x) \cdot E_K(y)$.
But it does not provide ciphertext indistinguishability (i.e., encryption is not randomized)

Additive Homomorphic Encryption


- Addition

- $E_K(m_1) \circ E_K(m_2) = E_K(m_1 + m_2)$

- Multiplication (by a constant c)

- $E_K(m)^c = E_K(m) \circ \dots \circ E_K(m) = E_K(c \cdot m)$

*“Bonus” operation
derived from addition*



Applications of PHE

- e-Voting
 - Calculate the total the votes without seeing plaintext votes
 - Protect the anonymity of the voters
- Digital cash
 - Ensure anonymity over financial transactions
- **Private Matching** / Private Set Intersection
 - Search for members of a watch list in an air flight passenger list



Threat Model (think about cloud computing)

1. Model in the status quo: Trusted

- Ask the cloud to do computation / search in plaintext

2. Honest-but-curious (aka semi-honest)

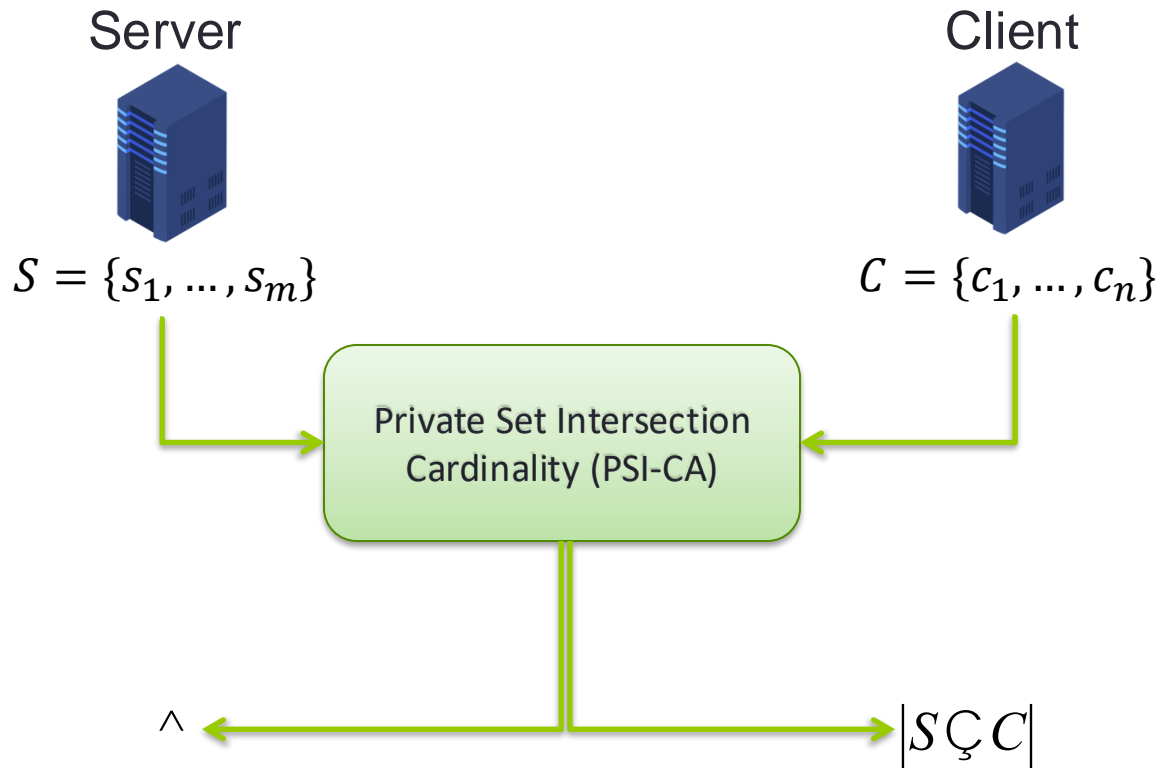
- Cloud cannot deviate from the protocol (i.e., honest)
- Cloud can try to learn more information; perform statistical inferences, or try to break the crypto (i.e., curious)
- Captures threats by curious system admins

3. Malicious

- Cloud can deviate arbitrarily from protocol

Private Set Intersection

Private Set Intersection Cardinality (PSI-CA)



Private Set Intersection

- Client has a set C of n items
- Server has a set S of m items
- We want to compute $C \cap S$ (or $|C \cap S|$) without revealing anything more about C and S

Approach:

1. Express C as a polynomial $P(X)$
2. Server evaluates $P(X)$ at each $s \in S$ using **additive** homomorphic encryption

Private Set Intersection

Client



Public key: K

Secret key: K'

$$C = \{c_1, \dots, c_n\}$$

$$P(X) = \prod_{i=1}^n (X - c_i) = \sum_{l=0}^n a_l X^l$$

$E_K(a_0), \dots, E_K(a_n)$



Server



Public key: K

$$S = \{s_1, \dots, s_m\}$$

For each $s_j \in S$:

- Pick a random r_j
- **Homomorphically evaluate** $P(s_j)$
- $E_K(r_j P(s_j) + s_j)$

Recall: Additive Homomorphic Encryption

- Addition
 - $E_K(m_1) \circ E_K(m_2) = E_K(m_1 + m_2)$
- Multiplication (by a constant c)
 - $E_K(m)^c = E_K(m) \circ \dots \circ E_K(m) = E_K(c \cdot m)$

Private Set Intersection

Client



$$C = \{c_1, \dots, c_n\}$$

Server



$$S = \{s_1, \dots, s_m\}$$

How does the **server** compute $E_K(r_j P(s_j) + s_j)$?

- Recall: For each s_j , pick a random r_j
- Evaluate $P(s_j)$ using $E_K(a_0), \dots, E_K(a_n)$ received from client
 - Recall that $P(X) = \prod_{i=1}^n (X - c_i) = \sum_{l=0}^n a_l X^l$
 - For $l = 0, \dots, n$:
 - compute s_j^l
 - then homomorphically compute $E_K(a_l) s_j^l = E_K(a_l s_j^l)$ (multiplication by a constant)
 - Homomorphically sum the terms by computing: $\prod_{l=0}^n E_K(a_l s_j^l) =$
 $E_K[\sum_{l=0}^n a_l s_j^l] = E_K[P(s_j)]$
 - $E_K[P(s_j)]^{r_j} \circ E_K[s_j] = E_K(r_j P(s_j) + s_j)$

Normal

addition

Homomorphic addition

Private Set Intersection

Client



Public key: K
Secret key: K'

$$C = \{c_1, \dots, c_n\}$$

$$P(X) = \prod_{i=1}^n (X - c_i) = \sum_{l=0}^n a_l X^l$$

$E_K(a_0), \dots, E_K(a_n)$



$E_K(r_1 P(s_1) + s_1), \dots, E_K(r_m P(s_m) + s_m)$



Server



Public key: K

$$S = \{s_1, \dots, s_m\}$$

For each $s_j \in S$:

- Pick a random r_j
- **Homomorphically evaluate** $P(s_j)$
- $E_K(r_j P(s_j) + s_j)$

- Client: perform intersection on the encrypted values:

- If $c_i = s_j$, then $P(s_j) = 0$, and thus $E_K(r_j P(s_j) + s_j) = E_K(s_j) = E_K(c_i)$
- Otherwise $E_K(r_j P(s_j) + s_j) = E_K(r)$, for some random r

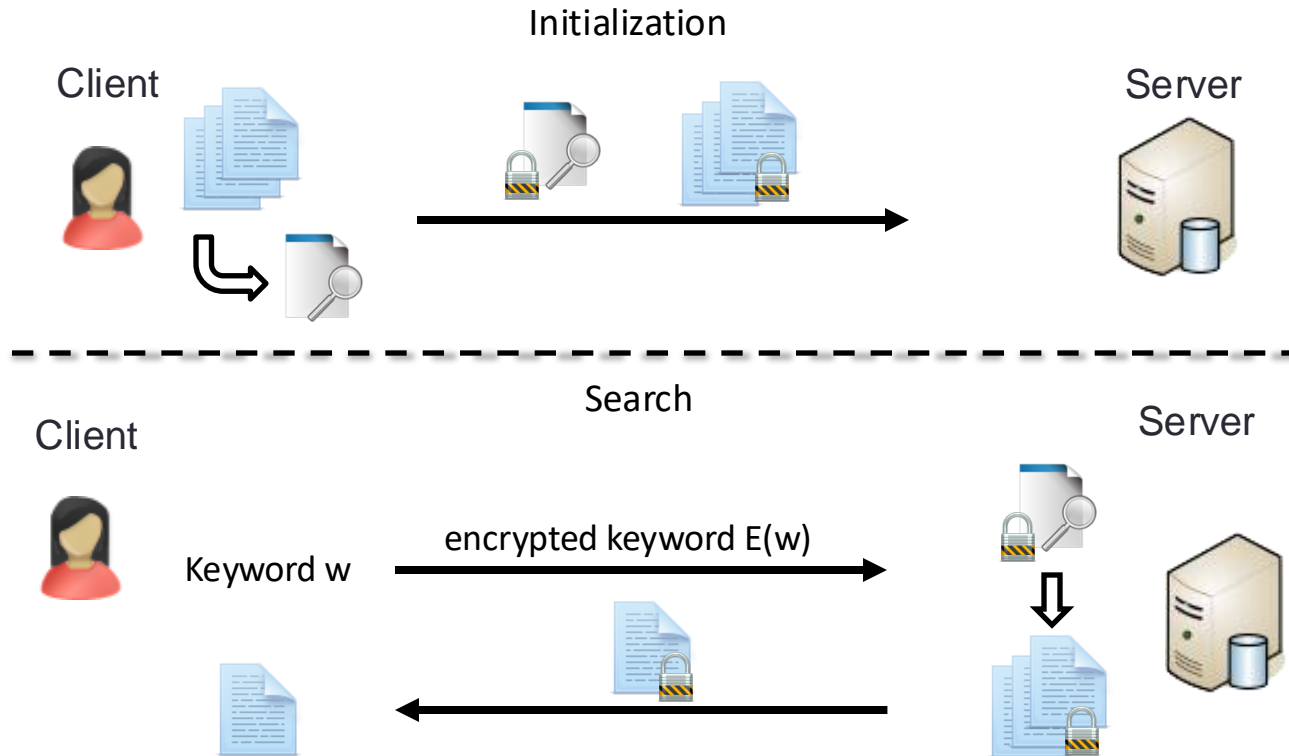
Searchable Encryption



Searchable Encryption

- Client wants to search for documents which contain a specific keyword
- Can the search be outsourced to a server without revealing the contents of the documents or the search keyword?
 - Client encrypts the documents, sends them to server
 - Client asks the server to return the (encrypted) documents containing a particular keyword

Searchable Encryption



Searchable Encryption

- Naive solution
 - Encrypt keywords (with a deterministic scheme)

Client



Search for
keyword w_2

$E(w_2)$



Server

Encrypted Keyword	Document IDs
$E(w_1)$	1, 7, 16
$E(w_2)$	3, 5
$E(w_3)$	7
$E(w_4)$	13, 11, 5, 2, 1

Cons: Index list will be HUGE!

Encrypted Index

Searchable Encryption

- Possible guarantees: the server learns only
 1. Keyword access pattern (i.e., last time this keyword was searched)
 2. Document access pattern (i.e., documents that are accessed for each keyword search)
- Reveals more in practice due to updates
 - e.g., add a document, delete a document

Access Pattern Leaks

- With auxiliary information:
 - Multi-user systems: correlate queries
 - Information about users who send the query: e.g., EMR of a patient is accessed by an oncologist
- Identify 80% of search queries on encrypted emails using access pattern alone
 - E.g., based on word distribution in emails

[IKK12] Islam, M., Mehmet Kuzu, and Murat Kantarcioglu. "Access pattern disclosure on searchable encryption: Ramification, attack and mitigation." NDSS 2012.

[CGPR15] Cash, D., Grubbs, P., Perry, J. and Ristenpart, T. "Leakage-abuse attacks against searchable encryption." ACM CCS 2015.

How to Make Accesses Oblivious?

“Doesn’t look like anything to me.”



Software Protection and ORAM (Extra Reading)

- [GO96] Oblivious RAM - Originally proposed for software protection by Goldreich and Ostrovsky
- Traditional approach to software protection:
 - Tamperproof CPU and encrypted program
 - Decryption key embedded in ROM inside CPU
 - For each instruction: fetch, decrypt, execute
 - Protect RAM content from the rest of the system
- RAM content can be encrypted, but program execution reveals the memory addresses accessed → motivation for **Oblivious RAM**

[GO96] Goldreich, and Ostrovsky. "Software protection and simulation on oblivious RAMs." Journal of the ACM (JACM) 1996.

References

- [RAD78] Rivest, Adleman, and Dertouzos. "On data banks and privacy homomorphisms." Foundations of secure computation 4.11 (1978).
- [CGPR15] Cash, D., Grubbs, P., Perry, J. and Ristenpart, T. "Leakage-abuse attacks against searchable encryption." ACM CCS 2015.
- [GO96] Goldreich, and Ostrovsky. "Software protection and simulation on oblivious RAMs." Journal of the ACM (JACM) 1996.
- [FNP04]: Freedman, Nissim, Pinkas. "Efficient private matching and set intersection." EUROCRYPT 2004.

Discussion Questions

- Why not just trust the cloud provider?
- What other problems could be solved using Private Set Intersection?
- Are there alternative architectures for searchable encryption?
 - Keep the index on the client?
 - Use two cloud providers?